

Fixed inference scenario

Frozen by the operator

Qwen2.5-7B-Instruct, BF16
2048 in / 128 out tokens
concurrency 32, vLLM 0.23.0
1 H100 80 GB, 700 W cap
P95 TTFT $\leq$ 1600 ms
P95 TPOT $\leq$ 20 ms
online budget 0.12 GPU-h

12-candidate serving grid

Max sequences

8, 16, 32

x

Max batched tokens

2048, 4096, 8192

+

Additional controls

2 no-chunk + 1 expert

Bounded evidence stages

CPU Sandbox

0 target GPU-h, screens all 12

Short Probe

1 H100, 64 requests, ranking

Full Verify

1 H100, 256 requests, recommendation

increasing fidelity and GPU-hour cost

A full sweep pays target-GPU time for every candidate; only full target-workload evidence can authorize a recommendation.