

Fixed inference scenario

Frozen by the operator

Qwen2.5-7B-Instruct, BF16
2048 in / 128 out tokens
concurrency 32, vLLM 0.23.0
1 x H100 80GB, 700 W cap
TTFT <= 1600 ms P95
TPOT <= 20 ms P95
online budget 0.12 GPU-h

12-candidate serving grid

Max sequences

8, 16, 32 (+ expert 256)

x

Max batched tokens

2048, 4096, 8192

x

Chunked prefill

enabled (+2 no-chunk controls)

Bounded evidence stages

CPU Sandbox (L0)

0 target GPU-h, screens all 12

Short Probe (L1)

1 H100, 64 requests, ranking

Full Verify (L4)

1 H100, 256 requests, release

increasing fidelity and GPU-hour cost

A full sweep pays target-GPU time for every candidate; only full target-workload evidence can authorize a recommendation.