

Scenario setup

Inference Scenario

model + workload trace
H100 node + SLO
GPU-hour budget

Scenario Compiler

typed search space
memory / support filter
12 immutable candidates

Bounded hybrid agent loop

LLM proposal + state guard

typed subgoal, guarded admission

IPIG selection

candidate x fidelity
cost-normalized gain

Energy Twin + tools

CPU projector, DCGM
Pareto + uncertainty

typed evidence + hashes

candidate + fidelity action

Verified release

Verified Output

verified Pareto set
confidence +
provenance, scripts

release only on L4

Bounded evidence path — TokenPower Sandbox

CPU Sandbox (L0)

0 target GPU-h
screens all 12 candidates

Short Probe (L1)

1 H100, 64 requests
100% frontier recall

Full Verify (L4)

1 H100, 256 requests
SLO-feasible release gate

 Evidence Store — telemetry, SHA-256 hashes, immutable provenance