

Scenario setup

Inference Scenario

model + workload trace
H100 node + SLO
GPU-hour budget

Scenario Compiler

typed search space
memory / support filter
12 immutable candidates



typed evidence + hashes

Bounded hybrid agent loop

LLM proposal + state guard

typed subgoal, guarded admission

IPIG selection

candidate + stage
cost-normalized gain

Energy Twin + tools

CPU projector, DCGM
Pareto + uncertainty

candidate + stage action

Verified recommendation

Verified Output

verified Pareto set
confidence +
provenance, scripts

recommend only after Verify

Bounded evidence path — TokenPowerSandbox

CPU Sandbox

0 target GPU-h
screens all 12 candidates

Short Probe

1 H100, 64 requests
rank + uncertainty update

Full Verify

1 H100, 256 requests
SLO-feasible recommendation gate



Evidence Store — telemetry, SHA-256 hashes, immutable provenance