

## Scenario setup

## Bounded hybrid agent loop

## Verified release

### Inference Scenario

model + workload trace  
H100 node + SLO  
GPU-hour budget

### Scenario Compiler

typed search space  
memory / support filter  
12 immutable candidates

### LLM proposal + state guard

typed subgoal, guarded admission

### IPIG selection

candidate x fidelity  
cost-normalized gain

### Energy Twin + tools

CPU projector, DCGM  
Pareto + uncertainty

### Verified Output

verified Pareto set  
confidence +  
provenance, scripts

*typed evidence + hashes*

*candidate + fidelity action*

*release only on L4*

## Bounded evidence path — TokenPowerSandbox

### CPU Sandbox (L0)

0 target GPU-h  
screens all 12 candidates

### Short Probe (L1)

1 H100, 64 requests  
100% frontier recall

### Full Verify (L4)

1 H100, 256 requests  
SLO-feasible release gate

Evidence Store — telemetry, SHA-256 hashes, immutable provenance